

# Fault-Tolerant Middleware: Communication

Reliable communication middleware for distributed systems




Kyle E. Benson

Department of Computer Science  
University of California, Irvine  
Irvine, California 92697  
`kebenson@uci.edu`

June 13, 2013

# Introduction

- ▶ Fault-tolerant communication goals
  - ▶ **Correctness** of messages, non-corruption guarantee
  - ▶ **Ordering** of messages
    - ▶ FIFO: If  $M_a$  sent before  $M_b$ ,  $M_a$  received before  $M_b$
    - ▶ Causal: If  $M_a$  causes  $M_b$  to be sent,  $M_a$  received before  $M_b$  at all processes
    - ▶ Total: If  $M_a$  delivered before  $M_b$  at process  $P_j$ ,  $M_a$  delivered before  $M_b$  at all other  $P_i$  too
  - ▶ **Delivery** guarantees, bounds on latency




# Foundations of Reliability

- ▶ How to make unicast reliable?
  - ▶ Prevent **omission failures**
    - ▶ Guarantee message delivery
    - ▶ Assume correct processes will deliver messages
    - ▶ Redeliver on *timeout*
    - ▶ No bound on time before reply
  - ▶ Guarantee ordering; ignore repeated messages
    - ▶ Sequence numbers
    - ▶ Timestamps
    - ▶ Logical clocks
  - ▶ Message integrity
    - ▶ Hashing
    - ▶ Certificates
    - ▶ Keys

# Reliable Group Communication

- ▶ Why not just use TCP?
  - ▶ Consider 100 machines each running 10 processes
  - ▶ 1000+ TCP connections at each machine
  - ▶ 1 million+ total!
  - ▶ **Not scalable!**
  - ▶ Relies on timeouts
  - ▶ Ordering harder
  - ▶ Similar problems with other client-server connection-oriented protocols
- ▶ How to exploit redundancy in communication paths?
- ▶ Answer: multicast trees




# Reliable Distributed Multicast

- ▶ **Observation:** distributed systems naturally address groups of processes
  - ▶ Coordinating events
  - ▶ Replica communication
  - ▶ Anycast
  - ▶ Reduction
  - ▶ Parallel computation
- ▶ Distributed process groups → multicast groups
  - ▶ IP multicast not always supported
  - ▶ Make application layer multicast
  - ▶ Let the *middleware* handle delivering message to proper groups
  - ▶ Decouples machine address from distributed function target
  - ▶ More efficient network usage

# Spread Toolkit

- ▶ Open-source tools for group communication
- ▶ [www.spread.org](http://www.spread.org)
- ▶ Hierarchical
  - ▶ Wide area: hop protocol
  - ▶ Local area: ring protocol
- ▶ Scales: tens of sites, tens of machines in each
- ▶ Bindings in many languages / platforms
  - ▶ C/C++
  - ▶ Java
  - ▶ Python, Perl, Ruby
  - ▶ Windows (98 - XP)
  - ▶ BSD / Linux / Solaris / Irix
  - ▶ Mac OS X

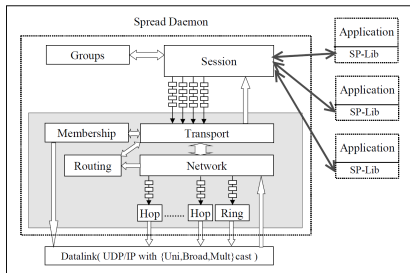


# Spread's Daemon-client Model

- ▶ Daemons provide messaging services
- ▶ User applications contact closest daemon for group communication
- ▶ Minimizes expensive membership changes
- ▶ Can tune number of daemons
- ▶ Wide area dissemination
  - ▶ Each site has one representative daemon for wide area dissemination
  - ▶ Routing trees rooted at each site
  - ▶ Each site = node in tree
  - ▶ Supports pruning, fast-retransmit, non-blocking send
- ▶ Daemons could even be run on routers
  - ▶ More for better performance
  - ▶ Fewer for less costly recovery

# Spread Architecture


- ▶ Several queues between session and transport
- ▶ Can support priorities
- ▶ Network module info → Routing module → Routing trees
- ▶ Can have multiple hops, only one ring
- ▶ Extended virtual synchrony
  - ▶ Handle network partitions
  - ▶ Handle re-merges
  - ▶ Joins
  - ▶ Leaves






# Spread's Hop Protocol

- ▶ Uses UDP/IP
- ▶ Handle losses hop-by-hop, not end-to-end
- ▶ Forward immediately, ignoring order
- ▶ NACKs for retransmit
  - ▶ NACK all outstanding packets
  - ▶ Wait timeout before requesting retransmit
  - ▶ Declares sender dead if retries > threshold
  - ▶ Latency bound
- ▶ cumulative ACKs
- ▶ Sliding window
- ▶ Token/leaky bucket for flow control




# Spread's Ring Protocol

- ▶ Multiple daemons in one site
- ▶ Local ordering, reliable dissemination
- ▶ Receives a token
  - ▶ Retransmits requested by previous holder
  - ▶ Receive messages
  - ▶ Send packets
  - ▶ Update and forward token



# Spread's Message Ordering

- ▶ Based on Lamport Time Stamp (logical clock)
- ▶ Sequence numbers
- ▶ Agreed
  - ▶ FIFO and Causal
  - ▶ Consistent across groups
- ▶ Safe
  - ▶ Consistent with agreed
  - ▶ Each site generates All-Received-Upto values
  - ▶ Disseminate across sites
  - ▶ Global All-Received-Upto values



# References

- ▶ Andreas Larsson. *Lab 2 Group Communication*.  
[http://www.cse.chalmers.se/edu/year/2011/course/TDA297/slides/presentation\\_lab2.ppt](http://www.cse.chalmers.se/edu/year/2011/course/TDA297/slides/presentation_lab2.ppt)
- ▶ ccsejc. *Distributed Systems*. <http://www.scribd.com/doc/23723915/28/Reliable-Communication>
- ▶ The Spread Toolkit.  
<http://www.spread.org/SpreadPlatforms.html>
- ▶ Prashant Shenoy. *CS 677 Remote Procedure Calls*.  
<http://lass.cs.umass.edu/~shenoy/courses/677/lectures.html>
- ▶ Yair Amir and Jonathan Stanton. *The Spread Wide Area Group Communication System*