

Fault-Tolerant Middleware: Communication

Reliable communication middleware for distributed systems

Kyle E. Benson

Department of Computer Science
University of California, Irvine
Irvine, California 92697

`kebenson@uci.edu`

May 30, 2013

Introduction


- ▶ Fault-tolerant communication goals
 - ▶ **Correctness** of messages, non-corruption guarantee
 - ▶ **Ordering** of messages
 - ▶ FIFO: If M_a sent before M_b , M_a received before M_b
 - ▶ Causal: If M_a causes M_b to be sent, M_a received before M_b at all processes
 - ▶ Total: If M_a delivered before M_b at process P_j , M_a delivered before M_b at all other P_i too
 - ▶ **Delivery** guarantees, bounds on latency

Foundations of Reliability

- ▶ How to make unicast reliable?
 - ▶ Prevent **omission failures**
 - ▶ Guarantee message delivery
 - ▶ Assume correct processes will deliver messages
 - ▶ Redeliver on *timeout*
 - ▶ No bound on time before reply
 - ▶ Guarantee ordering; ignore repeated messages
 - ▶ Sequence numbers
 - ▶ Timestamps
 - ▶ Logical clocks
 - ▶ Message integrity
 - ▶ Hashing
 - ▶ Certificates
 - ▶ Keys

Reliable Group Communication

- ▶ Why not just use TCP?
 - ▶ Consider 100 machines each running 10 processes
 - ▶ 1000+ TCP connections at each machine
 - ▶ 1 million+ total!
 - ▶ **Not scalable!**
 - ▶ Relies on timeouts
 - ▶ Ordering harder
 - ▶ Similar problems with other client-server connection-oriented protocols
- ▶ How to exploit redundancy in communication paths?
- ▶ Answer: multicast trees



Reliable Distributed Multicast

- ▶ **Observation:** distributed systems naturally address groups of processes
 - ▶ Coordinating events
 - ▶ Replica communication
 - ▶ Anycast
 - ▶ Reduction
 - ▶ Parallel computation
- ▶ Distributed process groups → multicast groups
 - ▶ IP multicast not always supported
 - ▶ Make application layer multicast
 - ▶ Let the *middleware* handle delivering message to proper groups
 - ▶ Decouples machine address from distributed function target