

Fault-Tolerant Middleware: Communication

Reliable communication middleware for distributed systems

Kyle E. Benson




Department of Computer Science
University of California, Irvine
Irvine, California 92697

`kebenson@uci.edu`

May 30, 2013

Introduction

- ▶ Fault-tolerant communication goals
 - ▶ **Correctness** of messages, non-corruption guarantee
 - ▶ **Ordering** of messages
 - ▶ FIFO: If M_a sent before M_b , M_a received before M_b
 - ▶ Causal: If M_a causes M_b to be sent, M_a received before M_b at all processes
 - ▶ Total: If M_a delivered before M_b at process P_j , M_a delivered before M_b at all other P_i too
 - ▶ **Delivery** guarantees, bounds on latency




Foundations of Reliability

- ▶ How to make unicast reliable?
 - ▶ Prevent **omission failures**
 - ▶ Guarantee message delivery
 - ▶ Assume correct processes will deliver messages
 - ▶ Redeliver on *timeout*
 - ▶ No bound on time before reply
 - ▶ Guarantee ordering; ignore repeated messages
 - ▶ Sequence numbers
 - ▶ Timestamps
 - ▶ Logical clocks
 - ▶ Message integrity
 - ▶ Hashing
 - ▶ Certificates
 - ▶ Keys

Reliable Group Communication

- ▶ Why not just use TCP?
 - ▶ Consider 100 machines each running 10 processes
 - ▶ 1000+ TCP connections at each machine
 - ▶ 1 million+ total!
 - ▶ **Not scalable!**
 - ▶ Relies on timeouts
 - ▶ Ordering harder
 - ▶ Similar problems with other client-server connection-oriented protocols
- ▶ How to exploit redundancy in communication paths?
- ▶ Answer: multicast trees




Reliable Distributed Multicast

- ▶ **Observation:** distributed systems naturally address groups of processes
 - ▶ Coordinating events
 - ▶ Replica communication
 - ▶ Anycast
 - ▶ Reduction
 - ▶ Parallel computation
- ▶ Distributed process groups → multicast groups
 - ▶ IP multicast not always supported
 - ▶ Make application layer multicast
 - ▶ Let the *middleware* handle delivering message to proper groups
 - ▶ Decouples machine address from distributed function target
 - ▶ More efficient network usage

Spread Toolkit

- ▶ Open-source tools for group communication
- ▶ www.spread.org
- ▶ Hierarchical
 - ▶ Wide area: hop protocol
 - ▶ Local area: ring protocol
- ▶ Scales: tens of sites, tens of machines in each
- ▶ Bindings in many languages / platforms
 - ▶ C/C++
 - ▶ Java
 - ▶ Python, Perl, Ruby
 - ▶ Windows (98 - XP)
 - ▶ BSD / Linux / Solaris / Irix
 - ▶ Mac OS X

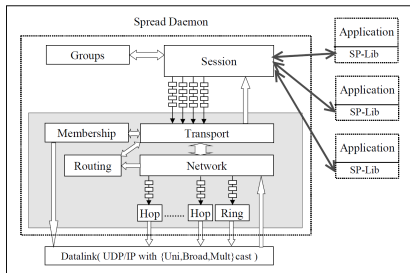


Spread's Daemon-client Model

- ▶ Daemons provide messaging services
- ▶ User applications contact closest daemon for group communication
- ▶ Minimizes expensive membership changes
- ▶ Can tune number of daemons
- ▶ Wide area dissemination
 - ▶ Each site has one representative daemon for wide area dissemination
 - ▶ Routing trees rooted at each site
 - ▶ Each site = node in tree
 - ▶ Supports pruning, fast-retransmit, non-blocking send
- ▶ Daemons could even be run on routers
 - ▶ More for better performance
 - ▶ Fewer for less costly recovery


Spread Architecture

- ▶ Several queues between session and transport
- ▶ Can support priorities
- ▶ Network module info → Routing module → Routing trees
- ▶ Can have multiple hops, only one ring
- ▶ Extended virtual synchrony
 - ▶ Handle network partitions
 - ▶ Handle re-merges
 - ▶ Joins
 - ▶ Leaves




Spread's Hop Protocol

- ▶ Uses UDP/IP
- ▶ Handle losses hop-by-hop, not end-to-end
- ▶ Forward immediately, ignoring order
- ▶ NACKs for retransmit
 - ▶ NACK all outstanding packets
 - ▶ Wait timeout before requesting retransmit
 - ▶ Declares sender dead if retries > threshold
 - ▶ Latency bound
- ▶ cumulative ACKs
- ▶ Sliding window
- ▶ Token/leaky bucket for flow control




Spread's Ring Protocol

- ▶ Multiple daemons in one site
- ▶ Local ordering, reliable dissemination
- ▶ Receives a token
 - ▶ Retransmits requested by previous holder
 - ▶ Receive messages
 - ▶ Send packets
 - ▶ Update and forward token



Spread's Message Ordering

- ▶ Based on Lamport Time Stamp (logical clock)
- ▶ Sequence numbers
- ▶ Agreed
 - ▶ FIFO and Causal
 - ▶ Consistent across groups
- ▶ Safe
 - ▶ Consistent with agreed
 - ▶ Each site generates All-Received-Upto values
 - ▶ Disseminate across sites
 - ▶ Global All-Received-Upto values



References

- ▶ Andreas Larsson. *Lab 2 Group Communication*.
http://www.cse.chalmers.se/edu/year/2011/course/TDA297/slides/presentation_lab2.ppt
- ▶ ccsejc. *Distributed Systems*. <http://www.scribd.com/doc/23723915/28/Reliable-Communication>
- ▶ The Spread Toolkit.
<http://www.spread.org/SpreadPlatforms.html>
- ▶ Prashant Shenoy. *CS 677 Remote Procedure Calls*.
<http://lass.cs.umass.edu/~shenoy/courses/677/lectures.html>
- ▶ Yair Amir and Jonathan Stanton. *The Spread Wide Area Group Communication System*